



Skeletal semantics

actors \leftrightarrow channels

“Skeletal Semantics”

Greetings

Name: Miguel Bartelsman

Project title: Skeletal semantics from actors to channels and back

Introduce theoretical backbone

First pillar: skeletal semantics

“Skeletal Semantics”

“First, let’s analyze what we mean by semantics”

Study of the rigorous meaning behind programming languages.

Given a syntactic construct in a language, its semantics tell us its result upon evaluation, correctness, and other properties.

Three types: we’ll focus on operational

Operational: repeatedly applying a set of reduction rules to a program

Skeletal Semantics

$$\frac{t_1 \rightarrow \text{true} \quad t_2 \rightarrow v_2}{\text{if } (t_1) \text{ then } (t_2) \text{ else } (t_3) \rightarrow v_2}$$

$$\frac{t_1 \rightarrow \text{false} \quad t_3 \rightarrow v_3}{\text{if } (t_1) \text{ then } (t_2) \text{ else } (t_3) \rightarrow v_3}$$

Here are the operational semantics rules for a hypothetical if-statement.

Top: guard is true

Bottom: guard is false

Skeletal Semantics

$$\frac{t_1 \rightarrow \text{true} \quad t_2 \rightarrow v_2}{\text{if } (t_1) \text{ then } (t_2) \text{ else } (t_3) \rightarrow v_2}$$

Premises

$$\frac{t_1 \rightarrow \text{false} \quad t_3 \rightarrow v_3}{\text{if } (t_1) \text{ then } (t_2) \text{ else } (t_3) \rightarrow v_3}$$

Conclusions

Evaluation relation

Explain the parts of a rule

Walk through one of them

“Skeletal Semantics”

Skeletal semantics is a theoretical framework for describing the operational semantics of a programming language.

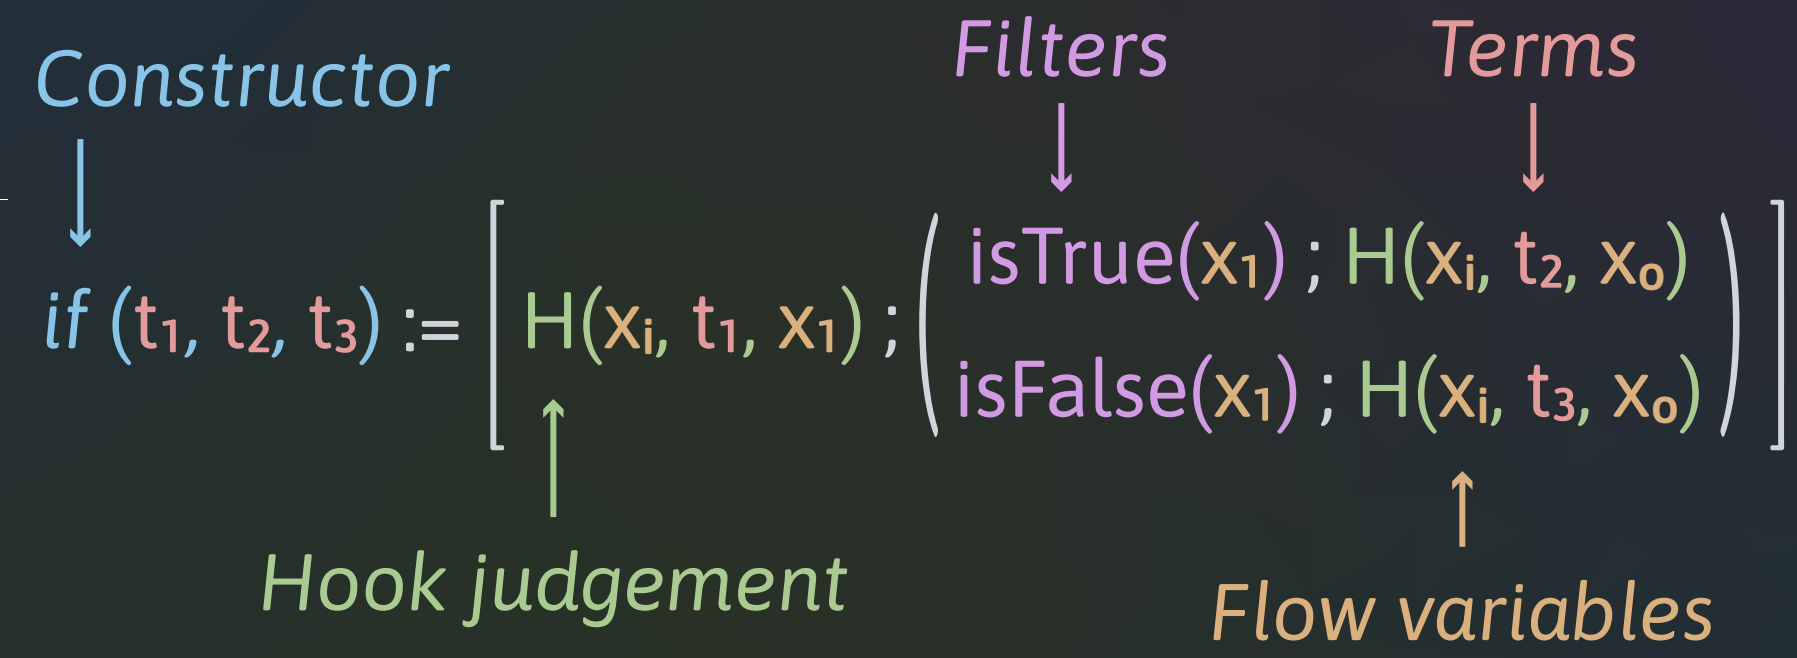
Many properties (shoutout to Laura)

Skeletal Semantics

$$if (t_1, t_2, t_3) := \left[H(x_i, t_1, x_1) ; \begin{pmatrix} \text{isTrue}(x_1) ; H(x_i, t_2, x_0) \\ \text{isFalse}(x_1) ; H(x_i, t_3, x_0) \end{pmatrix} \right]$$

This rule describes an if statement.

Skeletal Semantics



Made up of filters and hooks

Branching

Flow variables

Constructor

Properties of skeletal semantics:
structured and sequential

Introduce necro

Skel and Necro

```
val eval_if (xi, t) =  
  let If (t1, t2, t3) = t in  
  let x1 = eval (xi, t1) in  
  branch  
    let x1 = isTrue (x1) in eval (x1, t2)  
  or  
    let x1 = isFalse (x1) in eval (x1, t3)  
  end
```

Here is the same rule, written in the DSL
Skel

Walk through the parts

Skel and Necro

```
(* ... *)  
let eval_if =  
  function (xi, t) →  
  begin match expr with  
  | If (t1, t2, t3) →  
    let* x1 = apply1 eval (xi, t1) in  
    M.branch [  
      (function () →  
        let* x1 = apply1 isTrue x1 in  
        apply1 eval (x1, t2)  
      end) ;  
      (function () →  
        let* x1 = apply1 isFalse x1 in  
        apply1 eval (x1, t3)  
      end)  
    ]  
  | _ → M.fail ""  
  end  
(* ... *)
```

And here we have *a tiny part* of the
result of compiling the rule into OCaml.



Concurrency

The second pillar to the project is
concurrency

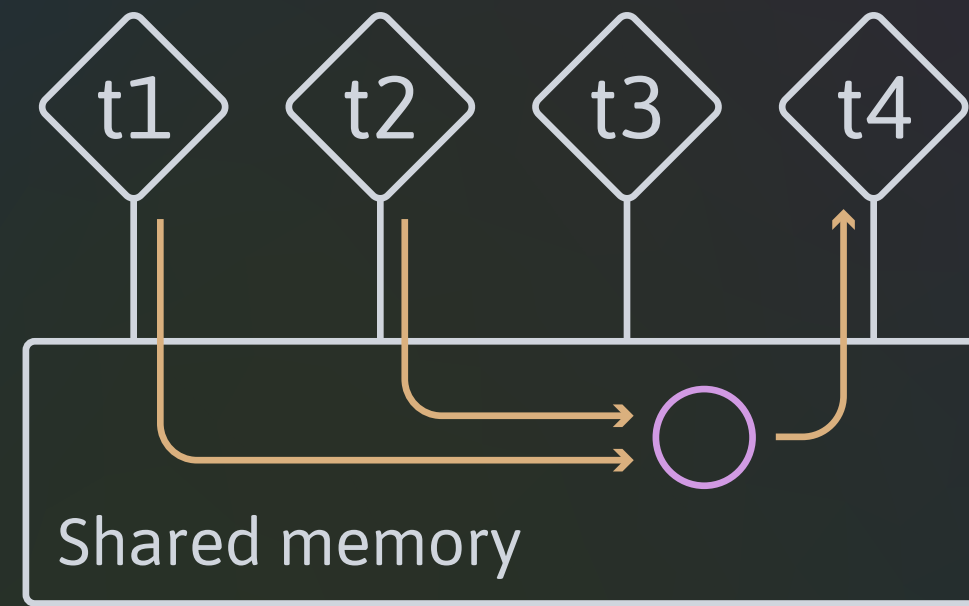
Shared memory
vs.
Message passing

There are two big approaches to concurrency

Shared memory, which is what most (all?) OSes use for concurrent execution of programs.

Message passing, which sees greater use in distributed computing.

Shared memory




Shared memory as the name implies

Processes or threads share one common memory heap where they store data.

Communication by writing to and reading from heap

Clashes, data races, lost data, poor scalability.

Message passing



Message passing addresses some of these issues.

Processes share data by sending messages. No memory is shared.

No lost data, no clashes, communication channels are abstractions that allow for extreme changes in scale for programs.


Actors
vs.
Channels

Withing MP concurrency there are two leading models

Actors : Erlang, Elixir, MPI

Channels : Go

Actors




Actors

Each process has an associated mailbox

Messages are sent directly to other processes and queues in their mailboxes

Channels



Channels

There are processes and channels

Messages can only be sent to channels and retrieved from channels. Making channels the interface for inter-process communications.

Equivalence between
actors and channels


These two models, asynchronous, are equivalent.

To show this, two translation schemes have been produced that show that each model can be represented in terms of the other.

Channels to actors

Channels

Actors



When translating channels to actors
Channels become actors executing a special program.

Processes become actors, instead of direct interfacing, they use messages to retrieve data from the channel

Actors to channels

Actors

Channels



Actors to channels,

Since processes have no mailbox, each actor is converted into a process-channel pair, where the channel acts as the associated mailbox of the process.

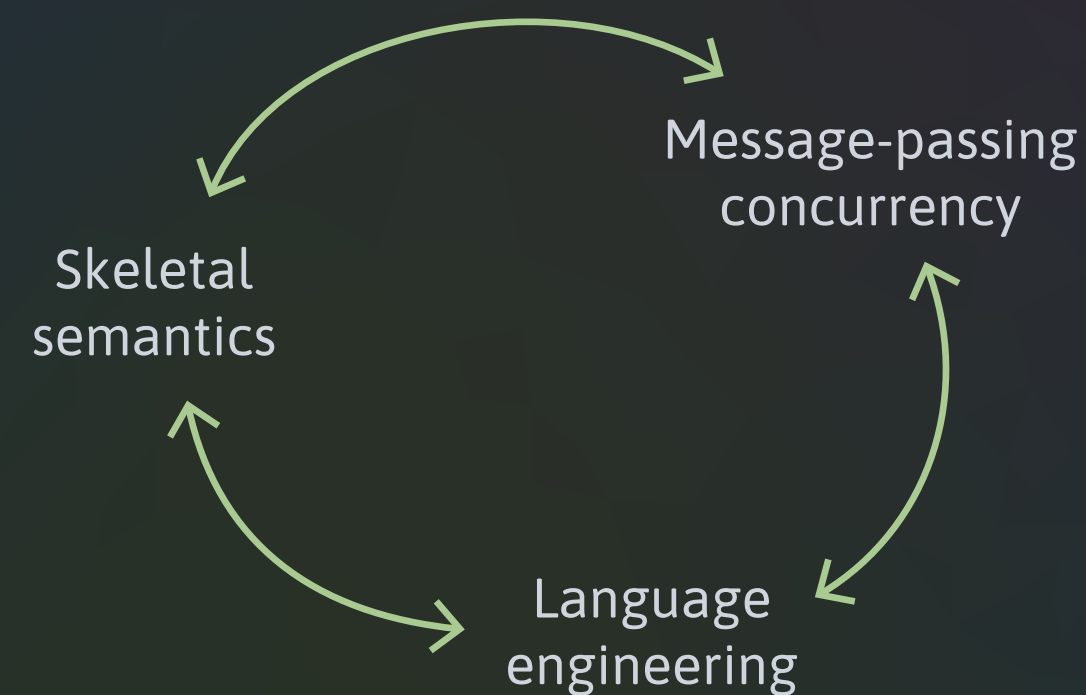
The project

Creation of a set of tools

viability of skeletal semantics and their associated tools

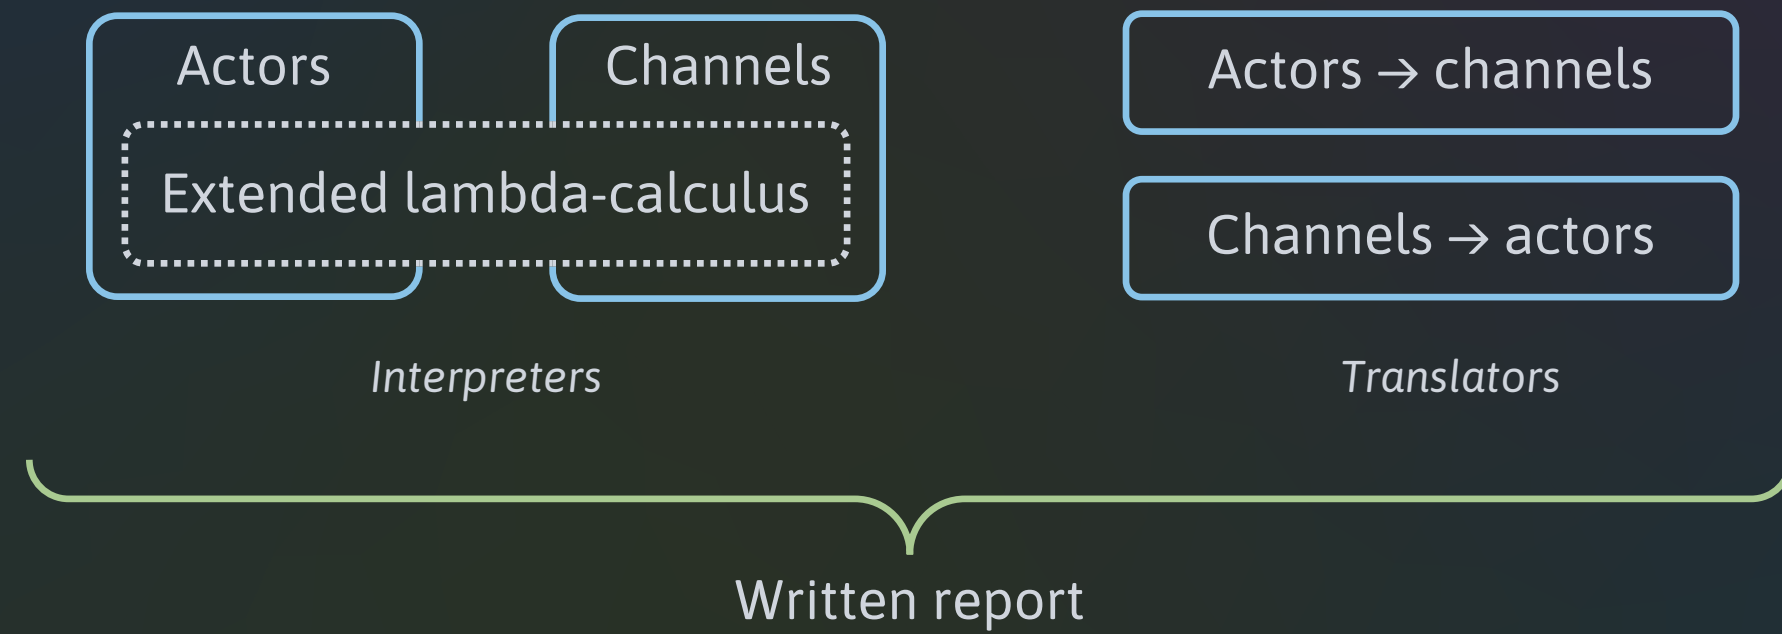
correct description and implementation of concurrent programming languages.

Goals



The project explores the space between skeletal semantics, mp concurrency, and language engineering.

Deliverables



To this end

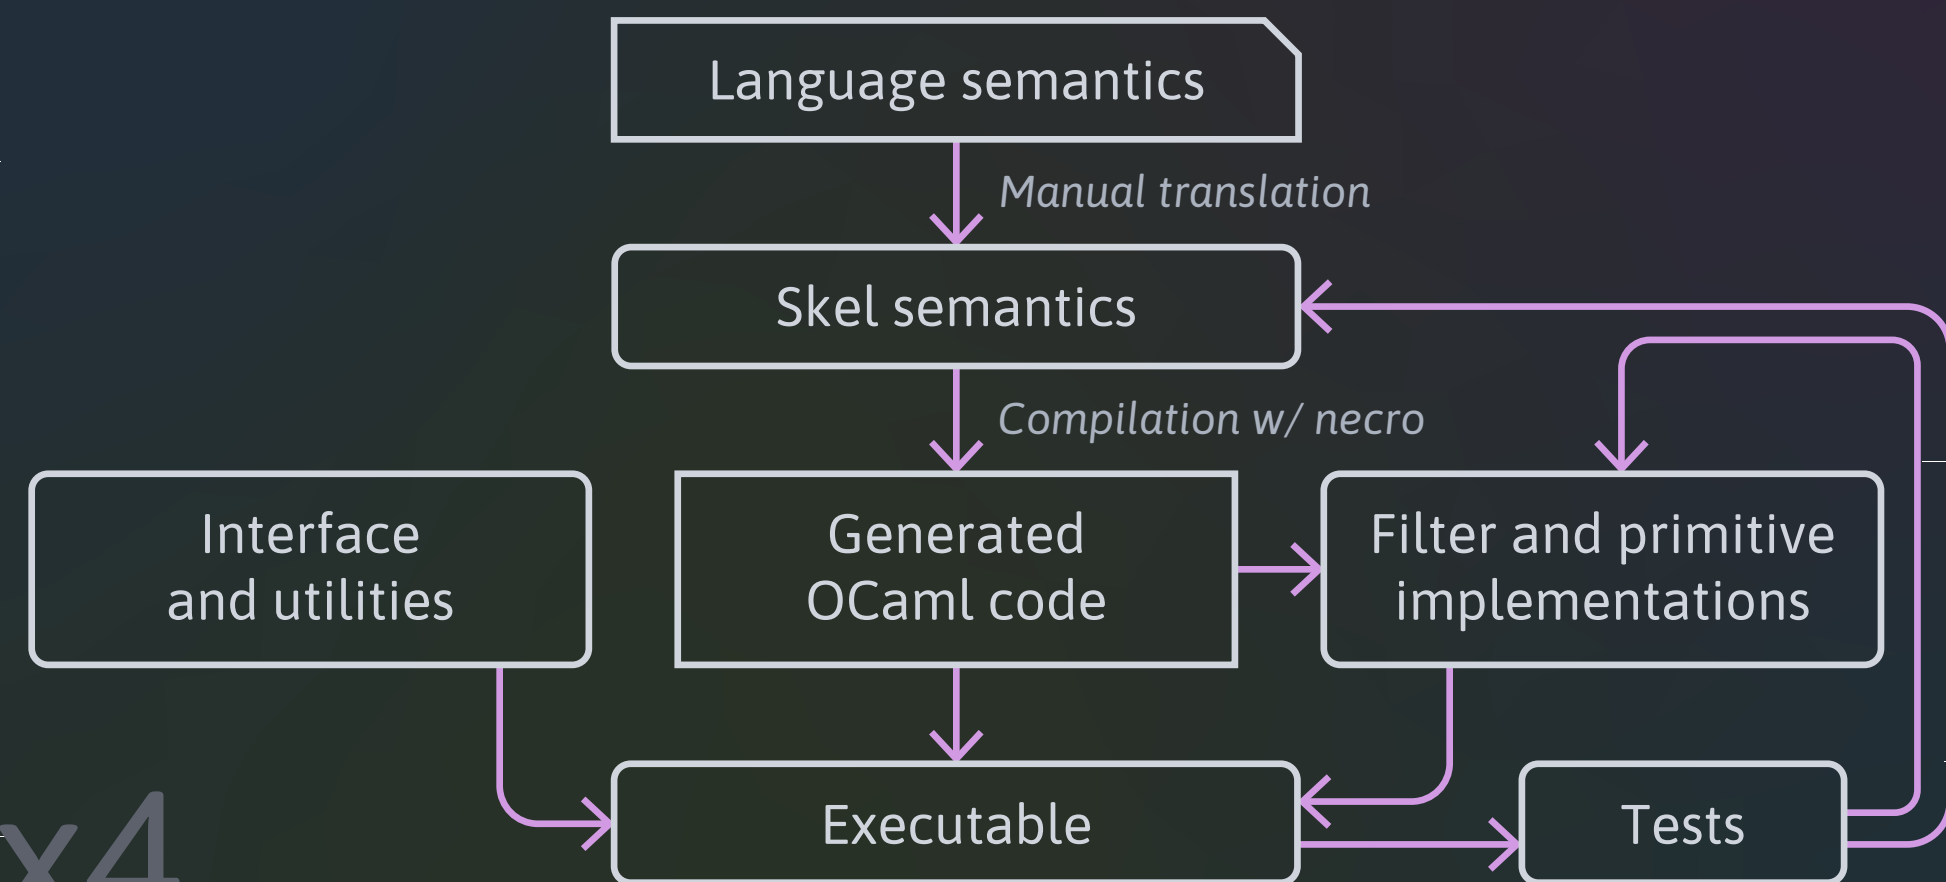
produce 4 software tools

Two interpreters for two concurrent languages built on top of an extended lambda calculus. One for each of the two models of concurrency describes.

Two translators, to go from one model to the other and back

And a written report that encompasses the entirety of the endeavour.

Workflow



x4

The workflow for the creation of the tools is as follows.

Existing semantics

Manually translated to skeletal sem.

Compiled to OCaml

Filters are treated as abstractions of the host language and implemented in OCaml, same for primitive types.

Interface functions are made

Wrapped together and tested


Results

Skeletal semantics are certainly powerful enough

But there's also plenty of issues to be considered.

Results

$$\begin{array}{l} P1 \parallel P2 \equiv P2 \parallel P1 \\ (P1 \parallel P2) \parallel P3 \equiv P1 \parallel (P2 \parallel P3) \end{array} \quad \frac{P1 \rightarrow P1'}{P1 \parallel P2 \rightarrow P1' \parallel P2}$$



For example, take parallel configurations.

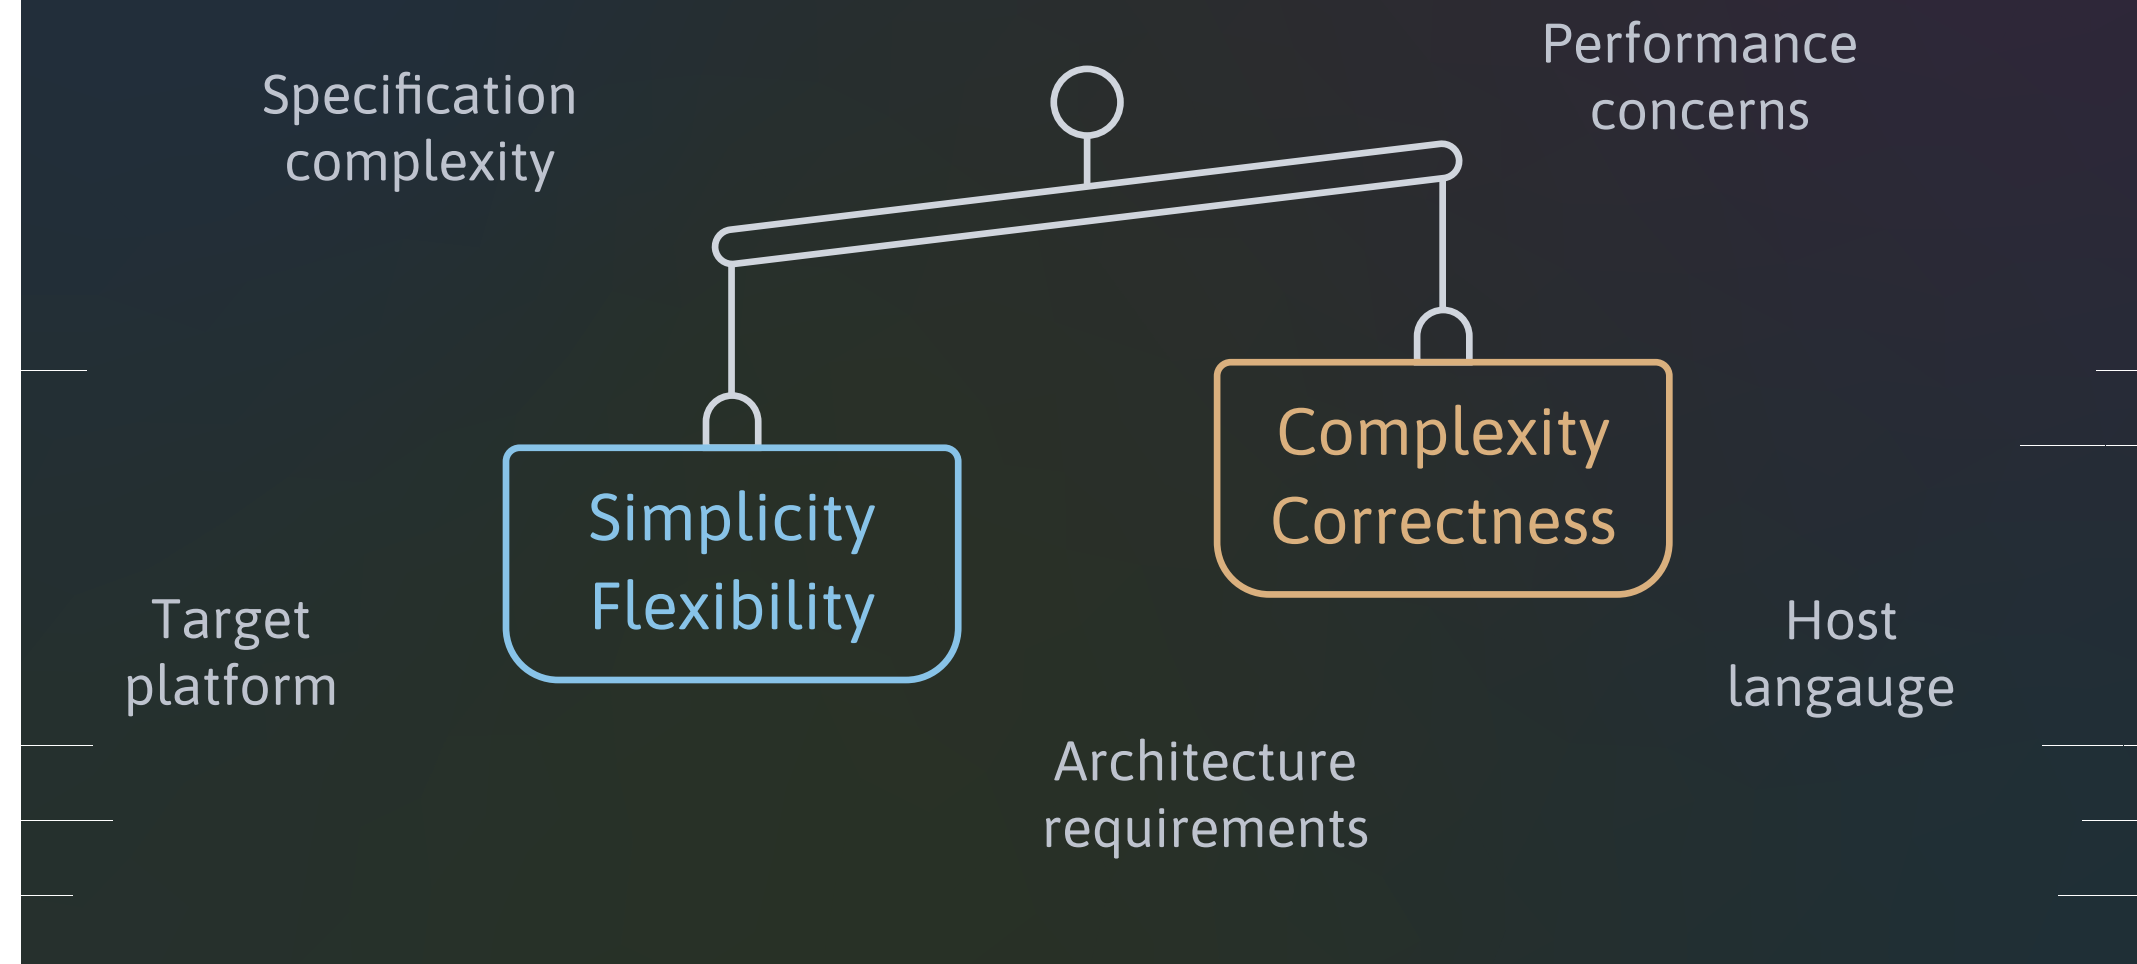
Commutative and distributive, that is, order and nesting do not matter when applying reduction rules

That is not the case when programming the semantics since computers do care about order and nesting.

The right abstractions and structures need to be chosen.

Such issues permeated the project as a whole

Results



Additionally, the use of skel sem is a balancing act.

Since filters allow for abstractions in the host language, it becomes a balancing act.

How much to abstract depends on many factors...